

RFC 001: Extend Quint with Row-Polymorphic Sum Types

Shon Feder

<2023-08-03 Thu>

Contents

1	Overview	2
1.1	Concrete syntax	2
1.1.1	Declaration	2
1.1.2	Constructors	3
1.1.3	Case analyses	3
1.2	Statics	4
1.2.1	Construction	4
1.2.2	Elimination	4
1.3	Dynamics	5
2	Discussion	5
2.1	Motivation	5
2.2	Context	6
2.2.1	Existing plans and previous work	6
2.2.2	The gist of extensible row-typed records and sum types	7
2.2.3	Quint’s current type system	7
2.3	Statics	7
2.3.1	Eliminating products and introducing sums	8
2.3.2	Introducing products and eliminating sums	11
2.4	Concrete Syntax	14
2.4.1	Why not support “polymorphic variants”	14
2.4.2	Declaration	15
2.4.3	Case analysis	21
2.4.4	Sketch of an alternative syntax	21
2.5	High-level implementation plan	23
2.6	Additional consideration	23

2.6.1	Pattern matching	23
2.6.2	User defined parametric type constructors	23
2.6.3	Drop the exotic operators	23

1 Overview

This section gives a concise overview of the principle proposal. In depth discussion follows in Discussion.

1.1 Concrete syntax

The concrete syntax falls into three categories, corresponding three aspects of the semantics:

- Declaring that a type is allowed in a context.
- Constructing an element of the declared type.
- Eliminating an element of the declared type (to derive an element of some other type).

Here we merely state the proposed syntax. See the discussion of concrete syntax for consideration of motivations and trade offs.

1.1.1 Declaration

The proposed syntax fits complements our syntax for records and follows the precedent of most languages that include support for general sum-types:

```
type T =
  | A(int)
  | B(str)
  | C
```

As indicated by `C`, labels with no associated expression type are allowed. These are sugar for a label associated with the unit type (i.e., the empty record). The above is therefore a representation of

```
type T =
  | A(int)
  | B(str)
  | C({})
```

1.1.2 Constructors

Internally, we need to add an operator `variant` similar to the “injection” operator from Leijen05:

```
<l = _> ::  $\tau$  .  $\rightarrow$  <l ::  $\tau$  |  $\tau$  > -- injection
```

This operation injects an expression of type τ into a sum-type that includes a variant associating label `l` with type τ . However, if we don’t want to expose the row-polymorphism to users, we’ll need a more restrictive typing. This is discussed in the section on typing rules.

We provide syntax sugar for users: when a sum type declaration is parsed, constructor operations for each variant will be generated internally. E.g., for the type `T` defined above, we will add to the context the following operators:

```
pure def A(x:int): T = variant(A, x)
pure def B(x:str): T = variant(B, x)
pure def C(): T      = variant(C, {})
```

1.1.3 Case analyses

Expressions of a sum type can be eliminated using a case analysis via a `match` statement. Following Rust’s syntax, we write that as

```
match e {
  | A(a) => ...
  | B(b) => ...
  | C    => ...
}
```

This construct can either be a primitive or syntax sugar for a more primitive `decompose` operator, discussed below.

We could also consider a case analysis `{ A(a) => e1, B(b) => e2, C => e3 }` – where `e1, e2, e3 : S` – as syntax sugar for an anonymous operator of type `T => S`. `match` would then be syntax sugar for operator application. This follows Scala’s case blocks or OCaml’s `function` keyword and would allow idioms such as:

```
setOfTs.map({ A(a) => e1 | B(b) => e2 | C => e3 })
```

instead of requiring

```
setOfTs.map(x => match x { A(a) => e1 | B(b) => e2 | C => e3 })
```

1.2 Statics

These type rules assume we keep our current approach of using quint strings for labels. But see my argument for simplifying our approach under Drop the exotic operators. See the discussion in Statics below for a detailed explanation and analysis.

1.2.1 Construction

The typing rule for constructing a variant of a sum type:

$$\frac{\Gamma \vdash e : (t, c) \quad \Gamma \vdash 'l' : str \quad \text{definedType}(s) \quad \text{free}(v)}{\Gamma \vdash \text{variant}('l', e) : (s, c \wedge s \sim \langle l : t | v \rangle)}$$

This rule is substantially different from Leijen05's

```
<l = i> :: r .  → <l ::  | r > -- injection
```

because we have decided not to expose the underlying row-polymorphism for sum types at this point. This introduces the non-trivial complication of needing to introduce the typing context onto our judgments (this is what we gesture at with the side condition `definedType(s)`).

1.2.2 Elimination

The typing rule for eliminating a variant of a sum type via case analysis:

$$\frac{\Gamma \vdash e : (s, c) \quad \Gamma, x_1 \vdash e_1 : (t, c_1) \quad \dots \quad \Gamma, x_n \vdash e_n : (t, c_n) \quad \Gamma, \langle v \rangle \vdash e_{n+1} : (t, c_{n+1}) \quad \text{fresh}(v)}{\Gamma \vdash \text{match } e \{ i_1 : x_1 \Rightarrow e_1, \dots, i_n : x_n \Rightarrow e_n \} : (t, c \wedge c_1 \wedge \dots \wedge c_n \wedge c_{n+1} \wedge s \sim \langle i_1 : t_1, \dots, i_n : t_n | v \rangle)}$$

This gives a rule in our system that is sufficient to capture Leijen05's

```
(l _ ? _ : _) :: r . <l ::  | r> → ( → ) → (<r> → ) → -- decomposition
```

since we can define decomposition for any label L via

```
def decomposeL(e: (L(a) | r), f: a => b, default : r => b) =
  match e {
    | L(x) => f(x)
    | r    => default(r)
  }
```

However we can define `match` as syntax sugar for the `decompose` primitive if we prefer.

1.3 Dynamics

The dynamics in the simulator should be straightforward and is not discussed here. Translation to Apalache for symbolic execution in the model checker is also expected to be relatively straight forward, since Apalache has a very similar form of row-based sum typing.

The general rules for eager evaluation can be found in PFPL, section 11.2. Additional design work for this will be prepared if needed.

This concludes the tl;dr overview of the proposal. The remaining is an indepth (still v. rough in places, discussion).

2 Discussion

2.1 Motivation

Quint’s type system currently supports product types. Product types (i.e., records, with tuples as a special case where fields are indexed by an ordinal) let us specify *conjunctions* of data types in a way that is verifiable statically. This lets us describe more complex data structures in terms of values of specific types that **must** be packaged together. E.g., we might define a rectangle by its length and width and a triangle by the lengths of its three sides. Using Quint’s existing syntax for product types, we’d specify this as follows:

```
type Rectangle =  
  { l : int  
    , w : int }  
type Triangle =  
  { a : int  
    , b : int  
    , c : int }
```

Quint’s type system does not yet have the the dual construct, sum types (aka “variants”, “co-products”, or “tagged unions”). Sum types specify *disjunctions* of data types in a way that is verifiable statically. This lets us describe mutually exclusive alternatives between distinct data structures that **may** occur together and be treated uniformly in some context. E.g., we might wish to specify a datatype for shapes, so we can work with collections that include both rectangles and triangles. Using one of the proposed syntax option that will be motivated in the following, this could be specified as

```

type Shape =
  | Rect(rectangle)
  | Tri(triangle)

```

Having both product types and sum types (co-product types) gives us a simple and powerful idiom for specifying families of data structures:

- We describe *what must be given together* to form a product of the specified type, and so *what we may always make use of* by projection when we are given such a product.
- We describe *which alternatives may be supplied* to form a co-product of a specified type, and so *what we must be prepared to handle* during case analysis when we are given such a co-product.

E.g., a `rectangle` is defined by *both* a length *and* a width, packaged together, while a `shape` is defined *either* by a rectangle *or* a triangle. With these definitions established, we can then go on to form and reason about collections of shapes like `Set[shape]`, or define properties common to all shapes like `isEquilateral : shape => bool`¹.

2.2 Context

2.2.1 Existing plans and previous work

We have always planned to support co-products in quint: their utility is well known and widely appreciated by engineers with experience in modern programming languages. We introduced co-products to Apache in <https://github.com/informalsystems/apalache/milestone/60> for the same reasons. The design and implementation of the latter was worked out by ?? (????) based on the paper “Extensible Records with Scoped Labels”. At the core of this design is a simple use of row-polymorphism that enables both extensible variants and extensible records, giving us products and co-products in a one neat package. The quint type system was also developed using row-polymorphism following this design. As a result of this forethought, extension of quint’s type system and addition of syntax to support sum-types is expected to be relatively straightforward.

¹The expressive power of these simple constructs comes from the nice algebraic properties revealed when values of a type are treated as equal up-to isomorphism. See, e.g., <https://codewords.recurse.com/issues/three/algebra-and-calculus-of-algebraic-data-types>

2.2.2 The gist of extensible row-typed records and sum types

The core concept in the row-based approach we’ve opted for is the following: we can use the same construct, called a “row”, to represent the *conjoined* labeled fields of a product type and the *alternative* labeled choices of a sum type. That the row types are polymorphic lets us extend the products and sums using row variables.

E.g., given the row

$$i_1 : t_1 , \dots , i_n : i_n | v$$

with each t_k -typed field indexed by label i_k for $1 \leq k \leq n$ and the free row variable v , then

$$\{i_1 : t_1 , \dots , i_n : i_n | v\}$$

is an open record conjoining the fields, and

$$\langle i_1 : t_1 , \dots , i_n : i_n | v \rangle$$

is an open sum type presenting the fields as (mutually exclusive) alternatives. Both types are extensible by substituting v with another (possibly open row). To represent a closed row, we omit the trailing $|v$.

2.2.3 Quint’s current type system

The current type system supported by quint is based on a simplified version of the constraint-based system presented in “Complete and Decidable Type Inference for GADTs” augmented with extensible (currently, just) records based on “Extensible Records with Scoped Labels”. A wrinkle in this genealogy is that quint’s type system includes neither GADTs nor scoped labels (and even the extensibility supported for records is limited). Moreover, due to their respective foci, neither of the referenced papers includes a formalization the complete statics for product types or sum types, and while we have implemented support for product types in quint, we don’t have our typing rules recorded.

2.3 Statics

This section discusses the typing judgements that will allow us to statically verify correct introduction and elimination of expressions that are variants of a sum type. The following considerations have informed the structure in which the proposed statics are discussed:

- Since sum-types are dual to product types, I consider their complementary typing rules together: first I will present the relevant rule for product types, then propose the complementary rule for sum types. This should help maintain consistency between the two kinds of typing judgements and ensure our implementations of both harmonize.
- Since we don't have our existing product formation or elimination rules described separate from the implementation, transcribing them here can serve to juice our intuition, supplement our design documentation, and perhaps give opportunity for refinement.
- Since our homegrown type system has some idiosyncrasies that can obscure the essence of the constructs under discussion, I precede the exposition of each rule with a text-book example adapted from Practical Foundations for Programming Languages. This is only meant as a clarifying touchstone.

2.3.1 Eliminating products and introducing sums

The elimination of products via projection and the introduction of sums via injection are the simplest of the two pairs of rules.

1. Projection Here is a concrete example of projecting a value out of a record using our current syntax:

```
val r : {a:int} = {a:1}
val ex : int = r.a
// Or, using our exotic field operator, which is currently the normal form
val ex_ : int = r.field("a")
```

A textbook rule for eliminating an expression with a finite product types can be given as

$$\frac{\Gamma \vdash e : \{i_1 \hookrightarrow \tau_1, \dots, i_n \hookrightarrow \tau_n\} \quad (1 \leq k \leq n)}{\Gamma \vdash e.i_k : \tau_k}$$

Where i is drawn from a finite set of indexes used to label the components of the product (e.g., fields of a record or positions in a tuple) and $i_j \hookrightarrow \tau_j$ maps the index i_j to the corresponding type τ_j .

This rule tells us that, when an expression e with a product type is derivable from a context, we can eliminate it by projecting out of e

with an index i_k (included in the type), giving an expression of the type t_k corresponding to that index. If we're given a bunch of stuff packaged together we can take out just the one part we want.

In our current system, typechecking the projection of a value out of a record implements the following rule

$$\frac{\Gamma \vdash e : (r, c) \quad \Gamma \vdash 'l' : str \quad fresh(t)}{\Gamma \vdash field(e, 'l') : (t, c \wedge r \sim \{ l : t | tail_t \})}$$

where

- we use the judgement syntax established in ADR5, in which $\Gamma \vdash e : (t, c)$ means that, in the typing context Γ , expression e can be derived with type t under constraints c ,
- $fresh(t)$ is a side condition requiring the type variable t to be fresh in Γ ,
- $'l'$ is a string literal with the internal representation l ,
- c are the constraints derived for the type r ,
- $tail_t$ is a free row-variable constructed by prefixing the fresh variable t with “tail”,
- $\{ l : t | tail_t \}$ is the open row-based record type with field, l assigned type t and free row- left as a free variable,
- and $r \sim \{ l : t | tail_t \}$ is a unification constraint.

Comparing the textbook rule with the rule in our system helps make the particular qualities and idiosyncrasies of our system very clear.

The most critical difference w/r/t to the complexity of the typing rules derives from the fact that our system subordinates construction and elimination of records to the language level operator application rather than implementing it via a special constructs that work with product indexes (labels) directly. This is what necessitates the consideration of the string literal $'l'$ in our premise. In our rule for type checking record projections we “lift” quint expressions (string literals for records and ints for products) into product indexes.

The most salient difference is the use of unification constraints. This saves us having to “inspect” the record type to ensure the label is present and obtain its type. These are both accomplished instead via the unification of r with the minimal open record including the fresh

type t , which will end up holding the inferred type for the projected value iff the unification goes through. This feature of our type system is of special note for our aim of introducing sum-types: almost all the logic for ensuring the correctness of our typing judgements is delegated to the unification rules for the row-types that carry our fields for product type and sum types alike.

2. Injection Here is a concrete example of injecting a value into a sum type:

```
val n : int = 1
val ex : A(int) = A(1)
```

A textbook rule for eliminating an expression belonging to a finite product type can be given as

$$\frac{\Gamma \vdash e : \tau_k \quad (1 \leq k \leq n)}{\Gamma \vdash i_k \cdot e : \langle i_1 \hookrightarrow \tau_1, \dots, i_n \hookrightarrow \tau_n \rangle}$$

Where i is drawn from a finite set of indexes used to label the possible alternatives of the co-product and $i_j \hookrightarrow \tau_j$ maps the index i_j to the corresponding type τ_j . We use $\langle \dots \rangle$ to indicate the labeling is now disjunctive and $i_k \cdot e$ as the injection of e into the sum type using label i_k . Note the symmetry with complementary rule for projection out of a record: the only difference is that the (now disjunctive) row (resp. (now injected) expression) is swapped from premise to conclusion (resp. from conclusion to premise).

This rule tells us that, when an expression e with a type t_k is derivable from a context, we can include it as an alternative in our sum type by injecting it with the label i_k , giving an element of our sum type. If we're given a thing that has a type allowed by our alternatives, it can be included among our alternatives.

If we were following the row-based approach outlined in Leijen05, then the proposed rule in our system, formed by seeking the same symmetry w/r/t projection out from a product, would be:

$$\frac{\Gamma \vdash e : (t, c) \quad \Gamma \vdash 'l' : str \quad fresh(s)}{\Gamma \vdash variant('l', e) : (s, c \wedge s \sim \{ l : t | tail_s \})}$$

Comparing this with our current rule for projecting out of records, we see the same symmetry: the (now disjunctive) row type is synthesized instead of being taken from the context.

However, if we don't want to expose the row-polymorphism to users, we need a more constrained rule that will ensure the free row variable is not surfaced. We can address this by replacing the side condition requiring s to be free with a side condition requiring that there it be defined, and in our constraint check that we can unify that defined type with a row that contains the given label with the expected type and is otherwise open.

$$\frac{\Gamma \vdash e : (t, c) \quad \Gamma \vdash 'l' : \text{str} \quad \text{definedType}(s) \quad \text{free}(v)}{\Gamma \vdash \text{variant}('l', e) : (s, c \wedge s \sim \langle l : t | v \rangle)}$$

Igor has voiced a strong preference that we do not allow anonymous or row-polymorphic sum types, which is why the last rule is proposed. It does complicate our typing rules, as it requires we draw from the typing context.

2.3.2 Introducing products and eliminating sums

Forming expressions of product types by backing them into records and eliminating expressions of sum types by case analysis exhibit the same duality, tho they are a bit more complex.

1. Packing expressions into records Here is a concrete example of forming a record using our current syntax:

```
val n : int = 1
val s : str = "one"
val ex : {a : int, b : str} = {a : n, b : s}
// Or, using our exotic Rec operator, which is currently the normal form
val ex_ : {a : int, b : str} = Rec("a", n, "b", s)
```

A textbook introduction rule for finite products is given as

$$\frac{\Gamma \vdash e_1 : \tau_1 \quad \dots \quad \Gamma \vdash e_n : \tau_n}{\Gamma \vdash \{i_1 \hookrightarrow e_1, \dots, i_n \hookrightarrow e_n\} : \{i_1 \hookrightarrow \tau_1, \dots, i_n \hookrightarrow \tau_n\}}$$

This tells us that for any expressions $e_1 : \tau_1, \dots, e_n : \tau_n$ derivable from our context we can form a product that indexes those n expressions by

i_1, \dots, i_n distinct labels, and packages all data together in an expression of type $\{i_1 \hookrightarrow \tau_1, \dots, i_n \hookrightarrow \tau_n\}$. If we're given all the things of the needed types, we can conjoint them all together into one compound package.

The following rule describes our current implementation:

$$\frac{\Gamma \vdash ('i_1', e_1 : (t_1, c_1)) \quad \dots \quad \Gamma \vdash ('i_n', e_n : (t_n, c_n)) \quad \text{fresh}(s)}{\Gamma \vdash \text{Rec}('i_1', e_1, \dots, 'i_n', e_n) : (s, c_1 \wedge \dots \wedge c_n \wedge s \sim \{i_1 : t_1, \dots, i_n : t_n\})}$$

The requirement that our labels show up in the premise as quint strings paired with each element of the appropriate type is, again, an artifact of the exotic operator discussed later, as is the **Rec** operator in the conclusion that consumes these strings. Ignoring those details, this rule is quite similar to the textbook rule, except we use unification of the fresh variable s to propagate the type of the constructed record, and we have to do some bookkeeping with the constraints from each of the elements that will be packaged into the record.

2. Performing case analysis Here is a concrete example of case analysis to eliminate an expression belonging to a sum type using the proposed syntax variants:

```
val e : T = A(1)
def describeInt(n: int): str = if (n < 0) then "negative" else "positive"
val ex : str = match e {
  | A(x) => describeInt(x)
  | B(x) => x
}
```

A textbook rule for eliminating an expression that is a variant of a finite sum type can be given as

$$\frac{\Gamma \vdash e : \langle i_1 \hookrightarrow \tau_1, \dots, i_n \hookrightarrow \tau_n \rangle \quad \Gamma, x_1 : \tau_1 \vdash e_1 : \tau \quad \dots \quad \Gamma, x_n : \tau_n \vdash e_n : \tau}{\Gamma \vdash \text{match } e \{ i_1 \cdot x_1 \hookrightarrow e_1 \mid \dots \mid i_n \cdot x_n \hookrightarrow e_n \} : \tau}$$

Note the complementary symmetry compared with the textbook rule for product construction: product construction requires n expressions to conclude with a single record-type expression combining them all;

while sum type elimination requires a single sum-typed expression and n ways to convert each of the n alternatives of the sum type to conclude with a single expression of a type that does not need to appear in the sum type at all.

The proposed rule for quint's type system is given without an attempt to reproduce our practice of using quint strings. This can be added in later if needed:

$$\frac{\Gamma \vdash e : (s, c) \quad \Gamma, x_1 \vdash e_1 : (t, c_1) \quad \dots \quad \Gamma, x_n \vdash e_n : (t, c_n) \quad \Gamma, \langle v \rangle \vdash e_{n+1} : (t, c_{n+1}) \quad \text{fresh}(v)}{\Gamma \vdash \text{match } e \{ i_1 : x_1 \Rightarrow e_1, \dots, i_n : x_n \Rightarrow e_n \} : (t, c \wedge c_1 \wedge \dots \wedge c_n \wedge c_{n+1} \wedge s \sim \langle i_1 : t_1, \dots, i_n : t_n \rangle)}$$

Compared with quint's rule for product construction we see the same complementary symmetry. However, we also see a striking difference: there is no row variable occurring in the product construction, but the row variable plays an essential function in sum type elimination of row-based variants. Row types in records are useful for extension operations (i.e., which we don't support in quint currently) and for operators that work over some known fields but leave the rest of the record contents variable. But the core idea formalized in a product type is that the constructor *must* package all the specified things together so that the recipient *can* chose any thing; thus, when a record is constructed we must supply all the things and there is no room for variability in the row. For sum types, in contrast the constructor *can* supply any one thing (of a valid alternate type), and requires the recipient *must* be prepared to handle every possible alternative.

In the presence of row-polymorphis, however, the responsibility of the recipient is relaxed: the recipient can just handle a subset of the possible alternatives, and if the expression falls under a label they are not prepared to handle, they can pass the remaining responsibility on to another handler.

Here is a concrete example using the proposed syntax, note how we narrow the type of T:

```
type T = A | B;;
def f(e) = match e {
  | A => 1
  | B => 2
  | _ => 0
```

```
}
```

```
// f can be applied to a value of type T
let a : T = A
let ex : int = f(A)
```

```
// but since it has a fallback for an open row, it can also handle any other var
let foo = Foo
let ex_ : int = f(foo)
```

Here's the equivalent evaluated in OCaml as proof of concept:

```
utop #
type t = [`A | `B]
let f = function
  | `A -> 1
  | `B -> 2
  | _ -> 0
let ex = f `A, f `Foo
;;
type t = [ `A | `B ]
val f : [> `A | `B ] -> int = <fun>
val ex : int * int = (1, 0)
```

All the features just discussed that come from row-polymorphism will not be available since we are choosing to suppress the row-typing.

2.4 Concrete Syntax

2.4.1 Why not support “polymorphic variants”

Our sum type system is based on row-polymorphism. The only widely used language I've found that uses row-polymorphism for extensible sum types is OCaml (and the alternative surface syntax, ReScript/Reason). For examples of their syntax for this interesting construct, see

ReScript <https://rescript-lang.org/docs/manual/latest/polymorphic-variant>

OCaml <https://v2.ocaml.org/manual/polyvariant.html>

These very flexible types are powerful, but they introduce challenges to the syntax (and semantics) of programs. For example, supporting anonymous variant types requires a way of constructing variants without predefined constructors. Potential approaches to address this include:

- A special syntax that (ideally) mirrors the syntax of the type.
- A special lexical marker on the labels (what ReScripts and OCaml do), e.g., ``A 1` or `#a 1` instead of `A(1)`.

However, we have instead opted to hide the row-polymorphism, and not expose this.

2.4.2 Declaration

1. Why use the `|` syntax to separate alternatives?

- In programming `|` is widely used for disjunction:
 - regex
 - boolean “or”
 - bitwise “or”
 - alternatives in BNF
 - parallel execution on the pi-calculus
- Many modern languages with support for sum types (or the more general union types) use `|`, including
 - Python
 - TypeScript
 - (of course) the MLs, Haskell, F#, Elm, OCaml, etc.

2. Why not use `,` to separate alternatives?

We have discussed modeling our concrete syntax for sum type declarations on Rust. But without changes to other parts of our language, this would leave the concrete syntax for type declarations too similar to record type declarations.

This similarity is aggravated by the fact that we currently don’t enforce any case-based syntactic rules to differentiate identifiers used for operator names, variables, or module names, and we are currently planning to extend this same flexibility to variant labels, just as we do for record field labels. Thus, we could end up with a pair of declarations like:

```
type T = {
    A : int,
    b : str,
}
```

```
type S = {
    A(int),
    b(str),
}
```

We are not confident that the difference between `_:_` and `_(_)` will be enough to keep readers from confusing the two.

But the chance of mistaking a record and sum type declaration is actually compounding a worse possible confusion: the part of a sum-type record enclosed in brackets is syntactically indistinguishable from a block of operators applications.

Given we tend to read data structures from the outside in, we feel confident that we were going to avoid confusion by requiring declarations to use `|` to demarcate alternatives:

```
type T = {
    a : int,
    B : str,
}
```

```
type S =
    | A(int)
    | b(str)
```

The latter seems much clearer to our team, and if we reflect this syntax also in `match`, it will give another foothold to help readers gather meaning when skimming the code.

3. Could we just copy Rust exactly?

In Rust, sum types are declared like this:

```
enum T {
    A(i64),
    B(i64),
}
```



```
    c
}
```

If we just adopted this syntax directly without also changing our syntax for records, we would introduce

- Breaks with our current convention around type declarations and use of keywords.
- May mislead users to try injecting values into the type via Rust's `T:A(x)` syntax, which clashes with our current module syntax.
- This would move us closer to Rust but further from languages like TypeScript and Python.

Rust's syntax makes pretty good sense within the context of the rest of Rust's syntax. Here is an overview:

declaration

```
struct Pair {
    fst: u64,
    snd: String
}

enum Either {
    Left(u64),
    Right(String)
}
```

construction

```
let s = Either::Left(4)
let p = Pair{
    fst: 4,
    snd: "Two"
}
```

elimination

```

let two = p.snd
let four = s match {
    Either::Left(_) => 4,
    Either::Right(_) => 4
}

```

Note how the various syntactic elements work together to give consistent yet clearly differentiated forms of expression for dual constructs:

- The prefix keyword for declarations consistently (**enum** vs. **struct**).
- A unique constructor name is used consistently for forming a record or variant.
- Lexical rules add clear syntactic markers:
 - Data constructors must begin with a capital letter
 - Field names of a struct (and method names of a trait, and module names) must begin with a lowercase letter.
 - This ensures the syntax for module access and sum type construction are unambiguous: `mymod::foo(x)` vs `MySumType::Foo(x)`.
 - The caps/lower difference also helps reflect the duality between record fields and sum type alternatives.

Compare with the syntax proposed for quint this RFC:

declaration

```

type Pair = {
    fst: int,
    snd: str,
}

```

```

type Either =
    | left(int),
    | right(str)

```

construction

```

let s = left(4)

```

```
let p = {
  fst: 4,
  snd: "Two"
}
```

elimination

```
let two = p.snd
let four = s match {
  | left(_) => 4
  | right(_) => 4
}
```

Following the current quint syntax, we don't differentiate declaration with keywords or data construction with prefix use of type names. But we make up for this lost signal with the evident difference between `|` and `,.` This also compensates for the inability to differentiate based on capitalization.

Finally, by adopting a syntax that is very similar to C++-like languages (like Rust), we risk presenting false friends. There are numerous subtle differences between quint and Rust, and if we lull users into thinking the syntax is roughly the same, they are likely to be disappointed when they discover that, e.g.,

- Unlike Rust, conditions in `if` must be wrapped in `(...)`
- Unlike Rust, conditions must have an `else` branch
- Unlike Rust, `let` is not used for binding
- Unlike Rust, type parameters are surrounded in `[...]` rather than `<...>`
- Unlike Rust, operators are declared with `def`
- Unlike Rust, type variables must begin with a lower-case letter

In short, given how many ways we differ from Rust syntax already, adopting Rust's syntax for sum types would be confusing in the context of our current syntax and possibly lead to incorrect expectations.

4. What if we want to be more Rust-like

If we want to use a syntax for sum types that is closer to, or exactly the same as, Rust's, then we should make at least the following changes to the rest of our syntax to preserve harmony:

- Require a prefix name relating to a type when constructing data, e.g., `Foo {a: 1, b: 2}` for constructing a record of type `Foo`. (Note this would mean dropping support for anonymous records types.)
- Introduce a lexical distinction between capitalized identifiers, used for data constructors (and type constants), and uncapitalized identifiers, used for records field labels and operators.
- Use keywords consistently for type declarations.

Taking these changes into account, we could render the previous quint examples thus:

declaration

```
type Pair = struct {
  fst: u64,
  snd: String
}
```

```
type Either = enum {
  Left(u64),
  Right(String)
}
```

construction

```
let s = Left(4)
let p = Pair {
  fst: 4,
  snd: "Two"
}
```

elimination

```
let two = p.snd
let four = s match {
  Left(_) => 4,
  Right(_) => 4,
}
```

2.4.3 Case analysis

Ergonomic support for sum types requires eliminators, ideally in the form of case analysis by pattern matching.

The proposed syntax is close to Rust's pattern syntax, modulo swapping `|` for `,` to be consistent with the type declaration syntax.

Here's some example Rust for comparison

```
match x {  
  A    => println!("a"),  
  B    => println!("b"),  
  C(v) => println!("cv"),  
  _    => println!("anything"),  
}
```

The `match` is a close analogue to our existing `if` expressions, and the reuse of the `=>` hints at the connection between case elimination and anonymous operators. The comma separated alternatives enclosed in `{...}` follow the variadic boolean action operators, which seems fitting, since sum types are disjunction over data.

One question if we adopt some form of pattern-based case analysis is how far we generalize the construct. Do we support pattern matching on scalars like ints and symbols? Do we support pattern matching to deconstruct compound data such as records and lists? What about sets? Do we allow pattern expressions to serve as anonymous operator (like Scala)?

My guess is that in most cases the gains in expressivity of specs would justify the investment, but it is probably best to start with limiting support to defined sum types and seeing where we are after that.

Until we have pattern matching introduced, we should flag a parsing error if the deconstructor argument is not a free variable, and inform the user that full pattern matching isn't yet supported.

2.4.4 Sketch of an alternative syntax

The syntax being proposed is chosen because it is familiar to Rust programmers, and is deemed sufficient so long as we don't need to expose the underlying row polymorphism. However, it has the down-sides of being very similar to the syntax for records, which might lead to confusion. I've also considered a more distinctive alternative which is also more consistently complementary to our records syntax. This group of alternatives follows Leijen05:

1. Declaration

Reflecting the fact that both records and sum types are based on rows, we use the same pairing $(:)$ and enumerating $(,)$ syntax, but signal to move from a conjunctive to a disjunctive meaning of the row by changing the brackets:

```
type T =  
  < A : int  
    , B : str  
  >
```

2. Injection

Injection uses a syntax that is dual with projection on records: `.` projects values out of products and injects them into co-products:

```
val a : T = <A.1>
```

Since this option gives a syntactically unambiguous representation of variant formation, there is no need to generate special injector operator, and `<_._>` can be the normal form for injection.

Annotation of anonymous sum types is clear and unambiguous:

```
def f(n: int): <C:int, D:str | s> =  
  if (n >= 0) <C.n> else <D."negative">
```

Compare with the corresponding annotation for a record type:

```
def f(n: int): {C:int, D:str | s} =  
  if (n >= 0) {C:n, D:"positive"} else {C:n, B:"negative"}
```

3. Elimination

Finally, elimination uses a syntax that is dual to record construction, signaling the similarity thru use of the surrounding curly braces, and difference via the presence of the fat arrows (this syntax is similar to the one proposed):

```
match e {  
  A : a => ...,  
  B : b => ...  
}
```

2.5 High-level implementation plan

Add parsing and extension of the IR for the syntax <https://github.com/informalsystems/quint/pull/1092>

Add generation of constructor operators <https://github.com/informalsystems/quint/pull/1093>

Add rules for type checking <https://github.com/informalsystems/quint/issues/244>

Add support in the simulator <https://github.com/informalsystems/quint/issues/1033>

Add support for converting to Apache <https://github.com/informalsystems/quint/issues/1034>

2.6 Additional consideration

2.6.1 Pattern matching

We may want to consider support for pattern matching at some point. I suspect the `match` construct will make users want this, but then again perhaps quint is simple enough that we can do without this complication. We'd also have to consider carefully how this would work with conversion to the Apache IR.

2.6.2 User defined parametric type constructors

If we want to gain the most value from the addition of sum types we should allow users to define parametric types such as `Option` and `Either`. See <https://github.com/informalsystems/quint/issues/1073>.

2.6.3 Drop the exotic operators

- Remove the special product type operators `fieldNames`, `Rec`, `with`, `label`, and `index`, or add support for first-class labels As is, I think these are not worth the complexity and overhead.

Compare our rule with the projection operation from “Extensible Records with Scoped Labels”, which does not receive the label ‘`l`’ as a string, instead treating it as a special piece of syntax:

$$(_ . l) :: r \rightarrow (l | r) \quad \{l :: _ \mid r\} \rightarrow _$$

Another point of comparison is Haskell’s “Datatypes with Field Labels”, which generates a projection function for each label, so that defining the datatype

```
data S = S1 { a :: Int, b :: String }
```

will produce functions

```
a :: S -> Int
b :: S -> String
```

1. Benefits

- (a) Simplified typing rules Abandoning this subordination to normal operator application would leave us with a rule like the following for record projection:

$$\frac{\Gamma \vdash e : (r, c) \quad \text{fresh}(t)}{\Gamma \vdash e.l : (t, c \wedge r \sim \{ l : t | \text{tail_}t \})}$$

This would allow us to remove the checks for string literals, instead leaving that to the outermost, syntactic level of our static analysis. A similar simplification would follow for record construction: the rule for `Rec` would not need to validate that it had received an even number of argument of alternating string literals and values, since this would be statically guaranteed by the parsing rules for the $\{l_1 : v_1, \dots, l_n : v_n\}$ syntax. This would be a case of opting for the “Parse, don’t validate” strategy.

- (b) Safer language

The added surface area introduced by these operators have contributed to several bugs, including at least the ones discussed here:

- <https://github.com/informalsystems/quint/issues/169>
- <https://github.com/informalsystems/quint/issues/816>
- <https://github.com/informalsystems/quint/issues/1081>

- (c) More maintainable code base

We are losing structure in our internal representation of expressions, which means we get less value out of typescript type system, have to do more work when converting to the Apache IR, and have to deal with more fussy details. By converting the rich structures of records into an internal representation `Rec(expr1, ..., expr1)`

- we lose the ability to ensure we are forming records correctly statically
- have to do manual arithmetic to ensure fields and values are paired up correctly
- essentially reduce ourselves to a “stringly typed” representation, relying entirely on detecting the `Rec` value in the `kind` field.